

Housing Price data exploration

July 23, 2017

This is the Exploration Data Analysis (EDA) part of kaggle competition - Housing price. The contents include: data overview and inspect, data cleaning and some feature engineering.

```
setwd('C:/Users/Bangda/Desktop/kaggle/housing-price')
library(tidyverse)
library(magrittr)
library(stringr)
library(e1071)
library(VIM)
library(mice)
library(gridExtra)
train <- read.csv('train.csv', header = TRUE, stringsAsFactors = FALSE)
test  <- read.csv('test.csv', header = TRUE, stringsAsFactors = FALSE)
dim(train)
```

```
## [1] 1460  81
```

```
dim(test)
```

```
## [1] 1459  80
```

We can see that there are many variables available to be the predictors, their meaning can be found here: <https://www.kaggle.com/c/house-prices-advanced-regression-techniques/data>.

1. Data overview

Since there are too many variables, we won't display the structure of `train` and `test` right here. We can have a basic idea of the dataset: many of variables are categorical variables, like types or conditions etc. In addition, the categories of categorical variables are not too much (unlike the situation in Titanic data sets). And we can see some variables may related together from the variable names, some variables seems like can be "grouped".

Let's check the data types of `train` and `test`.

```
# get the data type of each column
train_var_class <- sapply(train, class)
test_var_class  <- sapply(test, class)
table(train_var_class)
```

```
## train_var_class
## character  integer
##          43       38
```

```
table(test_var_class)
```

```
## test_var_class
## character  integer
##          43       37
```

```
# numeric variables
```

```
train_var_class[train_var_class %in% c('integer', 'numeric')] %>% names()
```

```
## [1] "Id" "MSSubClass" "LotFrontage" "LotArea"
```

```
## [5] "OverallQual" "OverallCond" "YearBuilt" "YearRemodAdd"
## [9] "MasVnrArea" "BsmtFinSF1" "BsmtFinSF2" "BsmtUnfSF"
## [13] "TotalBsmtSF" "X1stFlrSF" "X2ndFlrSF" "LowQualFinSF"
## [17] "GrLivArea" "BsmtFullBath" "BsmtHalfBath" "FullBath"
## [21] "HalfBath" "BedroomAbvGr" "KitchenAbvGr" "TotRmsAbvGrd"
## [25] "Fireplaces" "GarageYrBlt" "GarageCars" "GarageArea"
## [29] "WoodDeckSF" "OpenPorchSF" "EnclosedPorch" "X3SsnPorch"
## [33] "ScreenPorch" "PoolArea" "MiscVal" "MoSold"
## [37] "YrSold" "SalePrice"

# categorical variables
train_var_class[train_var_class %in% c('factor', 'character')] %>% names()
```

```
## [1] "MSZoning" "Street" "Alley" "LotShape"
## [5] "LandContour" "Utilities" "LotConfig" "LandSlope"
## [9] "Neighborhood" "Condition1" "Condition2" "BldgType"
## [13] "HouseStyle" "RoofStyle" "RoofMatl" "Exterior1st"
## [17] "Exterior2nd" "MasVnrType" "ExterQual" "ExterCond"
## [21] "Foundation" "BsmtQual" "BsmtCond" "BsmtExposure"
## [25] "BsmtFinType1" "BsmtFinType2" "Heating" "HeatingQC"
## [29] "CentralAir" "Electrical" "KitchenQual" "Functional"
## [33] "FireplaceQu" "GarageType" "GarageFinish" "GarageQual"
## [37] "GarageCond" "PavedDrive" "PoolQC" "Fence"
## [41] "MiscFeature" "SaleType" "SaleCondition"
```

But when we check the actual value of some “numerical” variables we can find that some values are kinds of rank or grade, they might be also considered as categorical variables.

2. Missing data

(1) Detect missing data

```
train_na_stat <- apply(train, 2, function(.data) sum(is.na(.data)))
train_na_stat[train_na_stat > 0]
```

```
## LotFrontage Alley MasVnrType MasVnrArea BsmtQual
## 259 1369 8 8 37
## BsmtCond BsmtExposure BsmtFinType1 BsmtFinType2 Electrical
## 37 38 37 38 1
## FireplaceQu GarageType GarageYrBlt GarageFinish GarageQual
## 690 81 81 81 81
## GarageCond PoolQC Fence MiscFeature
## 81 1453 1179 1406
```

```
test_na_stat <- apply(test, 2, function(.data) sum(is.na(.data)))
test_na_stat[test_na_stat > 0]
```

```
## MSZoning LotFrontage Alley Utilities Exterior1st
## 4 227 1352 2 1
## Exterior2nd MasVnrType MasVnrArea BsmtQual BsmtCond
## 1 16 15 44 45
## BsmtExposure BsmtFinType1 BsmtFinSF1 BsmtFinType2 BsmtFinSF2
## 44 42 1 42 1
## BsmtUnfSF TotalBsmtSF BsmtFullBath BsmtHalfBath KitchenQual
## 1 1 2 2 1
```

```
##      Functional  FireplaceQu  GarageType  GarageYrBlt  GarageFinish
##           2           730           76           78           78
##      GarageCars  GarageArea  GarageQual  GarageCond  PoolQC
##           1           1           78           78          1456
##           Fence  MiscFeature  SaleType
##          1169          1408           1
```

get the variables contain NA

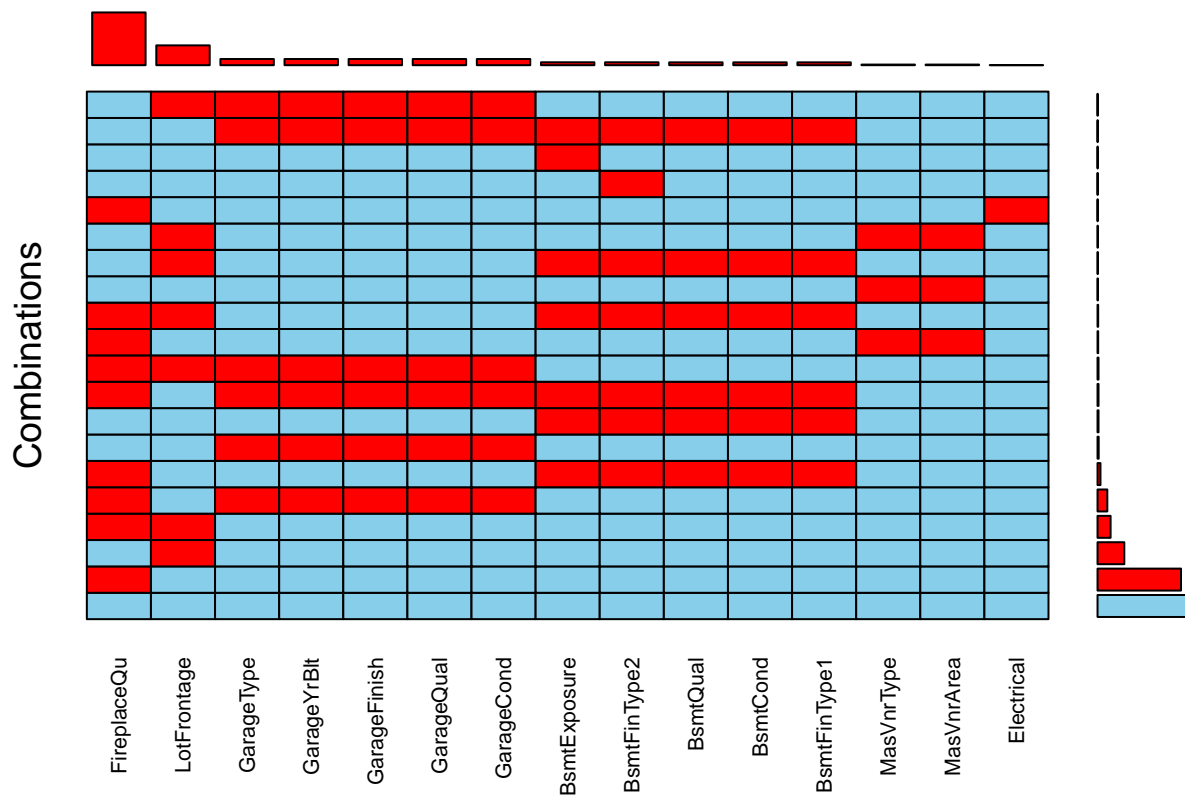
```
train_na_variable <- names(train_na_stat[train_na_stat > 0])
test_na_variable  <- names(test_na_stat[test_na_stat > 0])
```

The missing rates of some variables are more than 50%, we can remove them directly since the information loss are significant. Therefore we remove *Alley*, *PoolQC*, *Fence* *MiscFeature* from the data set.

```
train %<>%
  select(-Alley, -PoolQC, -Fence, -MiscFeature)
test %<>%
  select(-Alley, -PoolQC, -Fence, -MiscFeature)
# update
train_na_stat <- apply(train, 2, function(.data) sum(is.na(.data)))
test_na_stat  <- apply(test, 2, function(.data) sum(is.na(.data)))
train_na_variable <- names(train_na_stat[train_na_stat > 0])
test_na_variable  <- names(test_na_stat[test_na_stat > 0])
```

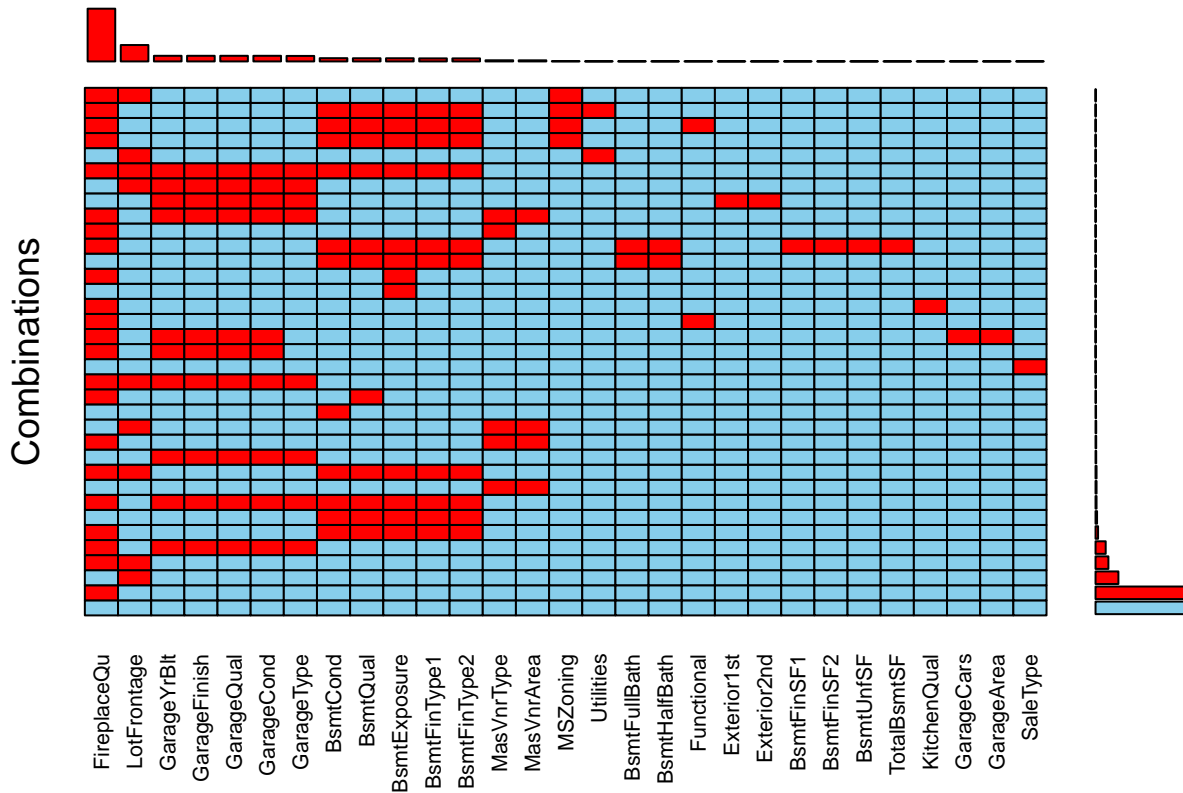
Then we visualize the missing data with the help of `aggr()` function in VIM.

```
train[, colnames(train) %in% names(train_na_stat[train_na_stat > 0])] %>%
  aggr(prop = FALSE, combined = TRUE, sortVars = TRUE, cex.axis = .7)
```



```
##
## Variables sorted by number of missings:
## Variable Count
## FireplaceQu 690
## LotFrontage 259
## GarageType 81
## GarageYrBlt 81
## GarageFinish 81
## GarageQual 81
## GarageCond 81
## BsmtExposure 38
## BsmtFinType2 38
## BsmtQual 37
## BsmtCond 37
## BsmtFinType1 37
## MasVnrType 8
## MasVnrArea 8
## Electrical 1
```

```
test[, colnames(test) %in% names(test_na_stat[test_na_stat > 0])] %>%
  aggr(prop = FALSE, combined = TRUE, sortVars = TRUE, cex.axis = .7)
```



```
##
## Variables sorted by number of missings:
## Variable Count
## FireplaceQu 730
## LotFrontage 227
```

```
## GarageYrBlt      78
## GarageFinish     78
## GarageQual       78
## GarageCond       78
## GarageType       76
## BsmtCond         45
## BsmtQual         44
## BsmtExposure     44
## BsmtFinType1     42
## BsmtFinType2     42
## MasVnrType       16
## MasVnrArea       15
## MSZoning         4
## Utilities        2
## BsmtFullBath     2
## BsmtHalfBath     2
## Functional       2
## Exterior1st      1
## Exterior2nd      1
## BsmtFinSF1       1
## BsmtFinSF2       1
## BsmtUnfSF        1
## TotalBsmtSF      1
## KitchenQual      1
## GarageCars       1
## GarageArea       1
## SaleType         1
```

We can also go with the “observations” side, we will count the missing variables for each observations.

```
train %>%
  apply(1, function(.row) sum(is.na(.row))) %>%
  sort(x = ., decreasing = TRUE) %>%
  '[' (1:10)
```

```
## [1] 11 11 11 11 11 11 11 10 7 7 7
```

```
test %>%
  apply(1, function(.row) sum(is.na(.row))) %>%
  sort(x = ., decreasing = TRUE) %>%
  '[' (1:10)
```

```
## [1] 12 12 11 11 11 11 11 11 8 8
```

So we know that the largest number of missing variables for observations is 11 in **train** and 12 in **test**. One of the practical rule is we can delete the observation if it has many missing variables. However, since the number of variables is much more greater than missing variables, we will not apply this rule in this case.

If we put variables contain NA along with the variable type we are suprised to find that almost all of them are categorical variables.

(2) Fill in missing data

One naive way to fill missing data for categorical data is using mode. So we will start from that and go back here later for more complicated methods like classification algorithms (decision trees) etc.

Since there is no built-in function to compute mode, we will define it by ourselves.

```
getMode <- function(x, na.rm = TRUE) {
  # get mode for character vector
  if (na.rm) {
    sort(table(x[!is.na(x)]), decreasing = TRUE)[1] %>% names()
  } else {
    sort(table(x), decreasing = TRUE)[1] %>% names()
  }
}
```

Gathering variables with NA,

```
train_na_set <- train[, colnames(train) %in% train_na_variable]
test_na_set <- test[, colnames(test) %in% test_na_variable]
```

Then apply our filling rule: use median to fill missing data for numeric variables, use mode to fill missing data for categorical variables,

```
fillNA <- function(x) {

  if (sum(is.na(x)) == 0) return(x)

  if (class(x) %in% c('integer', 'numeric')) {
    x[which(is.na(x))] <- median(x, na.rm = TRUE)
  } else {
    x[which(is.na(x))] <- getMode(x, na.rm = TRUE)
  }
  x
}
```

Test,

```
table(train_na_set$MasVnrType)
```

```
##
## BrkCmn BrkFace    None    Stone
##      15      445     864     128
```

```
train_na_set$MasVnrType %<>% fillNA()
table(train_na_set$MasVnrType)
```

```
##
## BrkCmn BrkFace    None    Stone
##      15      445     872     128
```

Finally we apply the function on data set, since using `apply()` with has some data type issues, we use for - loop here.

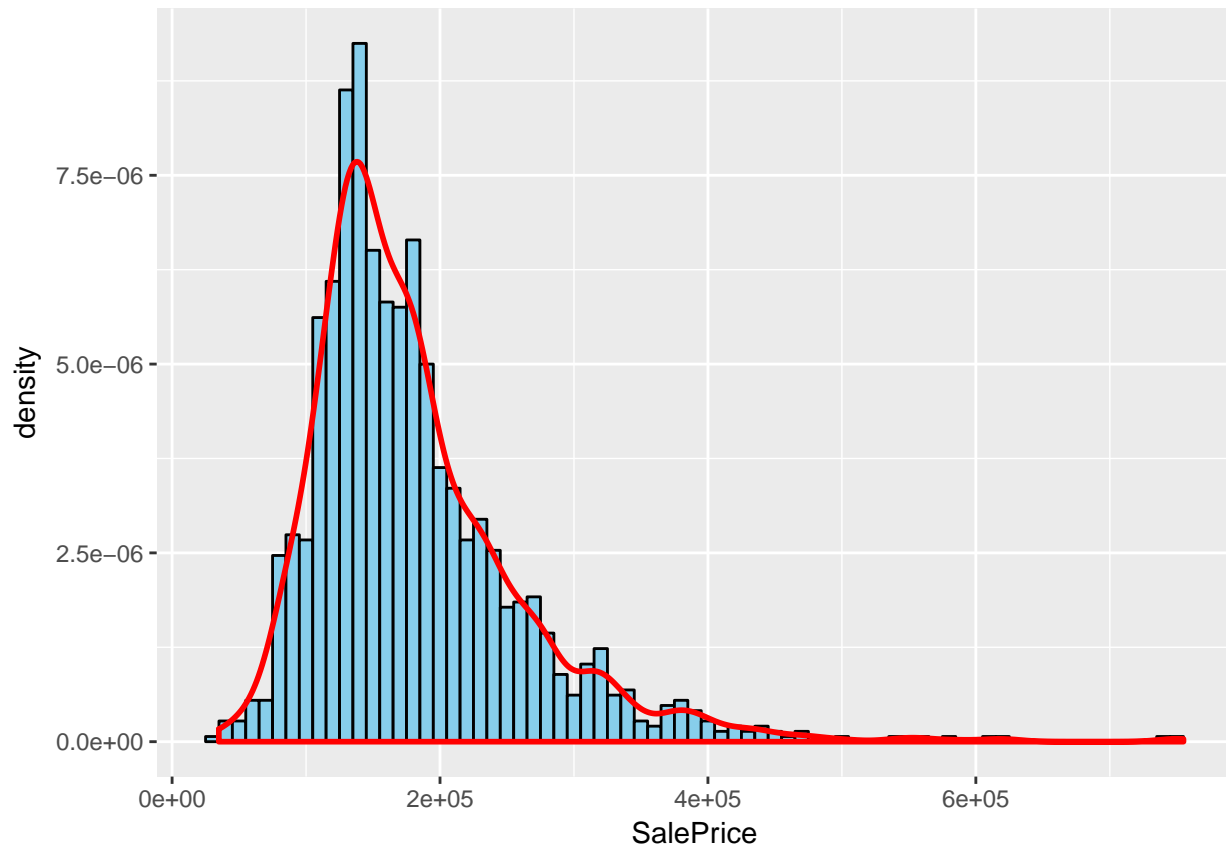
```
train_na_filled <- train
test_na_filled <- test
for (i in 1:ncol(train)) {
  train_na_filled[, i] <- fillNA(train[, i])
}
for (i in 1:ncol(test)) {
  test_na_filled[, i] <- fillNA(test[, i])
}
```

3. EDA and Feature Engineering

(1) Price

In order to satisfy the assumption of modeling (e.g. linear regression), we expect the *SalePrice* has normal distribution,

```
ggplot(train, aes(x = SalePrice)) +  
  geom_histogram(aes(y = ..density..), binwidth = 10000,  
    color = 'black', fill = 'skyblue') +  
  geom_density(aes(y = ..density..), size = 1, col = 'red')
```

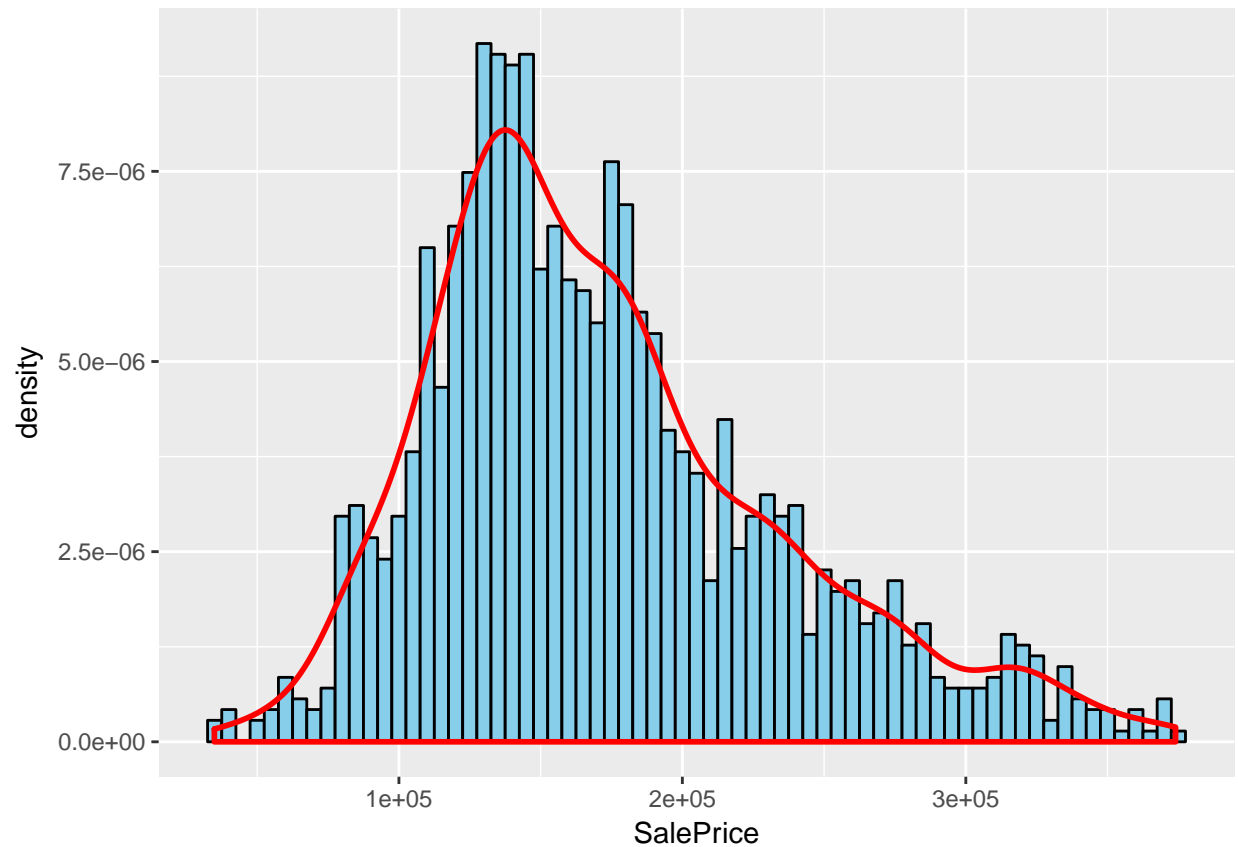


```
skewness(train$SalePrice)
```

```
## [1] 1.879009
```

The histogram says the distribution of *SalePrice* is not perfectly normal, it's a little right skewed since there are some relative high price.

```
train %>%  
  filter(SalePrice <= quantile(SalePrice, .97)) %>%  
  ggplot(aes(x = SalePrice)) +  
    geom_histogram(aes(y = ..density..), binwidth = 5000,  
      color = 'black', fill = 'skyblue') +  
    geom_density(aes(y = ..density..), size = 1, col = 'red')
```

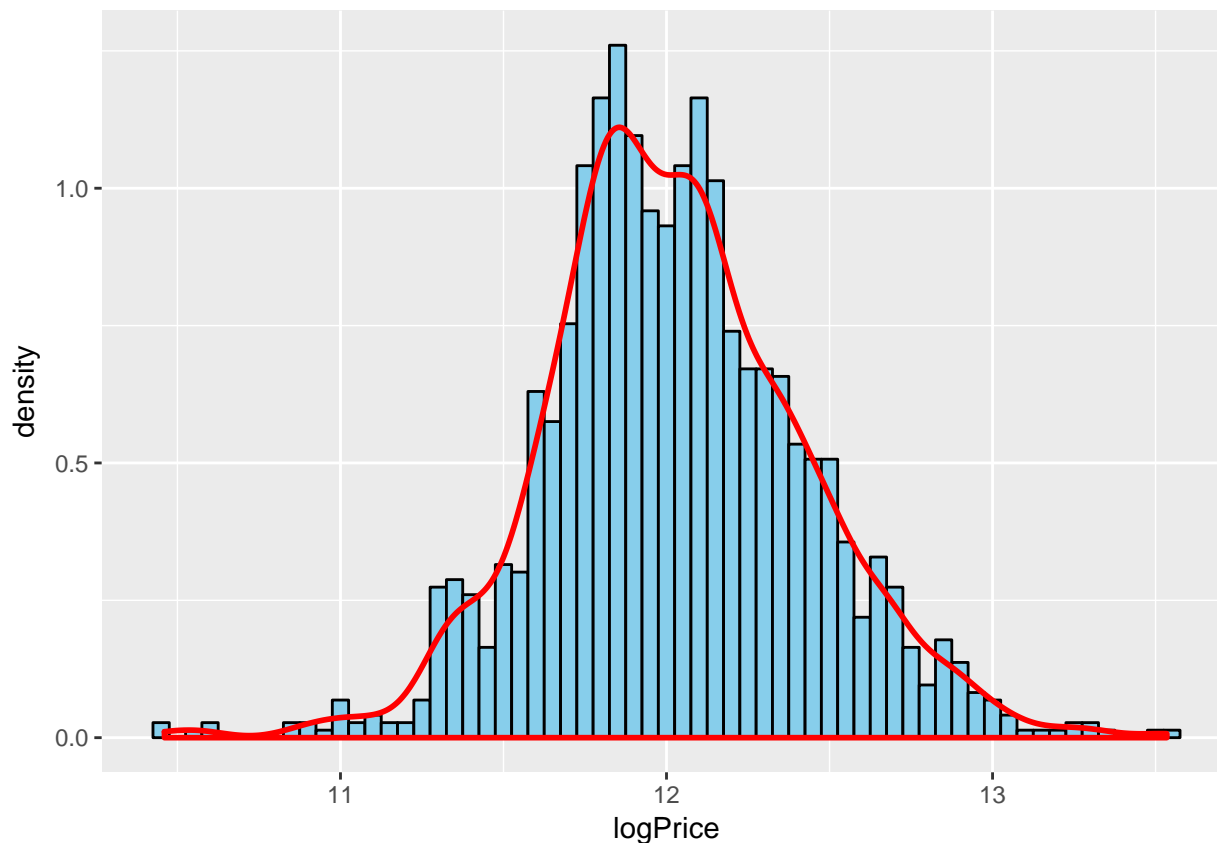


```
train %>%
  filter(SalePrice <= quantile(SalePrice, .97)) %>%
  skewness(SalePrice)
```

```
## [1] 0.7975751
```

Based on the evaluation of the result, we can also take the logarithm

```
train %>%
  mutate(logPrice = log(SalePrice)) %>%
  ggplot(aes(x = logPrice)) +
  geom_histogram(aes(y = ..density..), binwidth = 0.05,
                 color = 'black', fill = 'skyblue') +
  geom_density(aes(y = ..density..), size = 1, col = 'red')
```

Now it looks better. Hence we can draw a conclusion that if the distribution is skewed, we can try to use logarithm transformation to fix it since it can “scale” the extreme value. Also don’t forget to take exponential transformation to convert back.

```
train %<>% mutate(logSalePrice = log(SalePrice))
train_na_filled %<>% mutate(logSalePrice = log(SalePrice))
skewness(train$logSalePrice)
```

```
## [1] 0.1210859
```

(2) Variable cluster

We observed that some variables have same prefix, they are considered to describe different attributes of same object. For example, variables with prefix *Bsmt* are descriptions of basement. Our goal is try to find those groups.

```
# get the first 4 characters of names
prefix <- names(train) %>% str_sub(1, 3) %>% unique()
prefix_ref <- data.frame(prefix, stringsAsFactors = FALSE)
# count the variables with same prefix
prefix_ref$count <- sapply(prefix_ref$prefix,
                           function(pattern) str_detect(names(train), pattern) %>% sum())
# filter the count with at least 3
prefix_ref %<>%
  filter(count >= 3)
# get the variables that could have groups
grouped_var_idx <- sapply(
```

```

str_extract_all(paste0(names(train)),
                paste(prefix_ref$prefix, collapse = '|')),
function(num) length(num) > 0)
names(train)[grouped_var_idx]

```

```

## [1] "LotFrontage" "LotArea"      "LotShape"      "LandContour"
## [5] "LotConfig"   "Condition1"    "Condition2"    "OverallCond"
## [9] "Exterior1st" "Exterior2nd"   "ExterQual"     "ExterCond"
## [13] "BsmtQual"    "BsmtCond"      "BsmtExposure"  "BsmtFinType1"
## [17] "BsmtFinSF1"  "BsmtFinType2"  "BsmtFinSF2"    "BsmtUnfSF"
## [21] "TotalBsmtSF" "BsmtFullBath"  "BsmtHalfBath"  "GarageType"
## [25] "GarageYrBlt" "GarageFinish"  "GarageCars"    "GarageArea"
## [29] "GarageQual"  "GarageCond"    "SaleType"      "SaleCondition"
## [33] "SalePrice"   "logSalePrice"

```

Actually, we can check the descriptions of variables, and roughly classify them into 4 categories: (1) Interior, like *LotArea*, *Utilities*; (2) Exterior, like *MSSubClass*, *LotFrontage*; (3) Location and transporations, like *street*, *Condition1*; (4) Other attributes, like *YearBuilt*, *YearRemodAdd*.

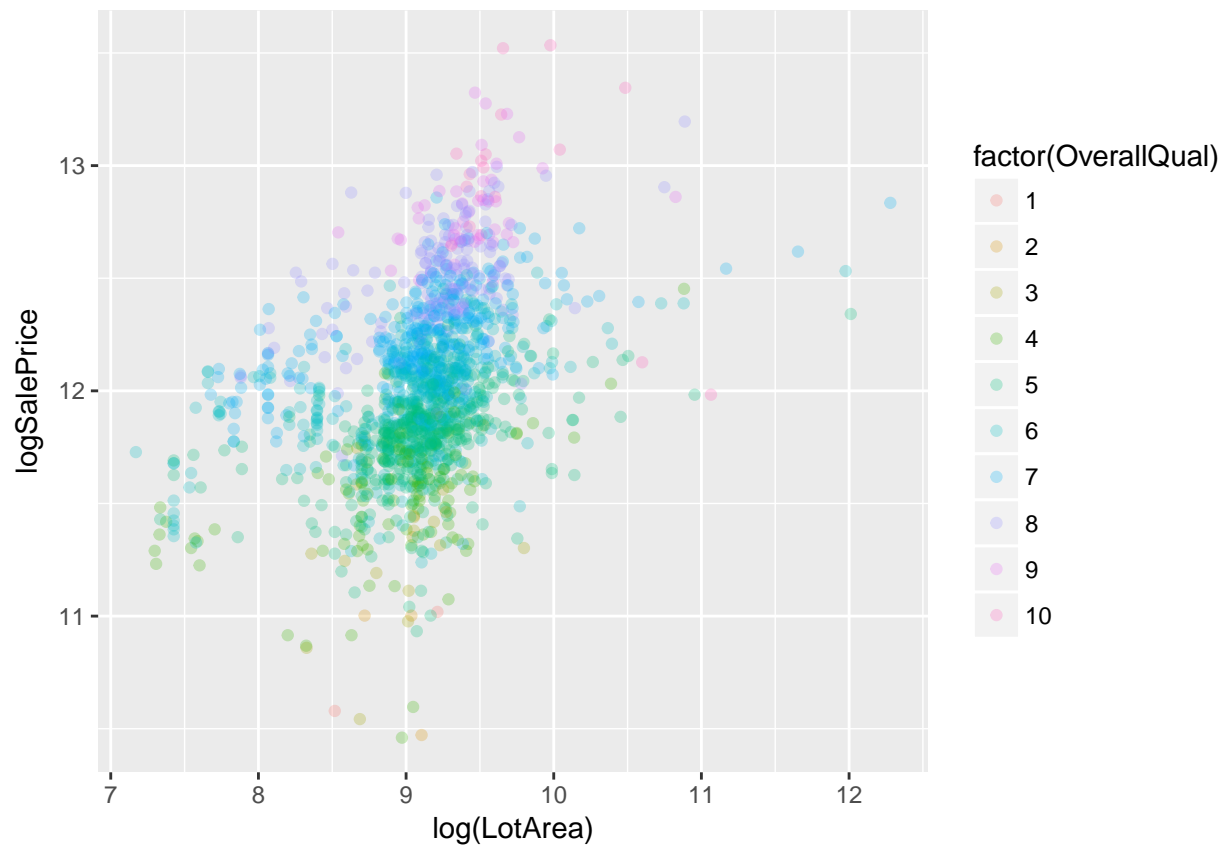
(3) Interior

Typically *LotArea* is a key factor that everyone will consider when buying a house, we can see that there is a linear positive correlation between *LotArea* and *logSalePrice*. Adding *OverallQual* makes this statement more reasonable.

```

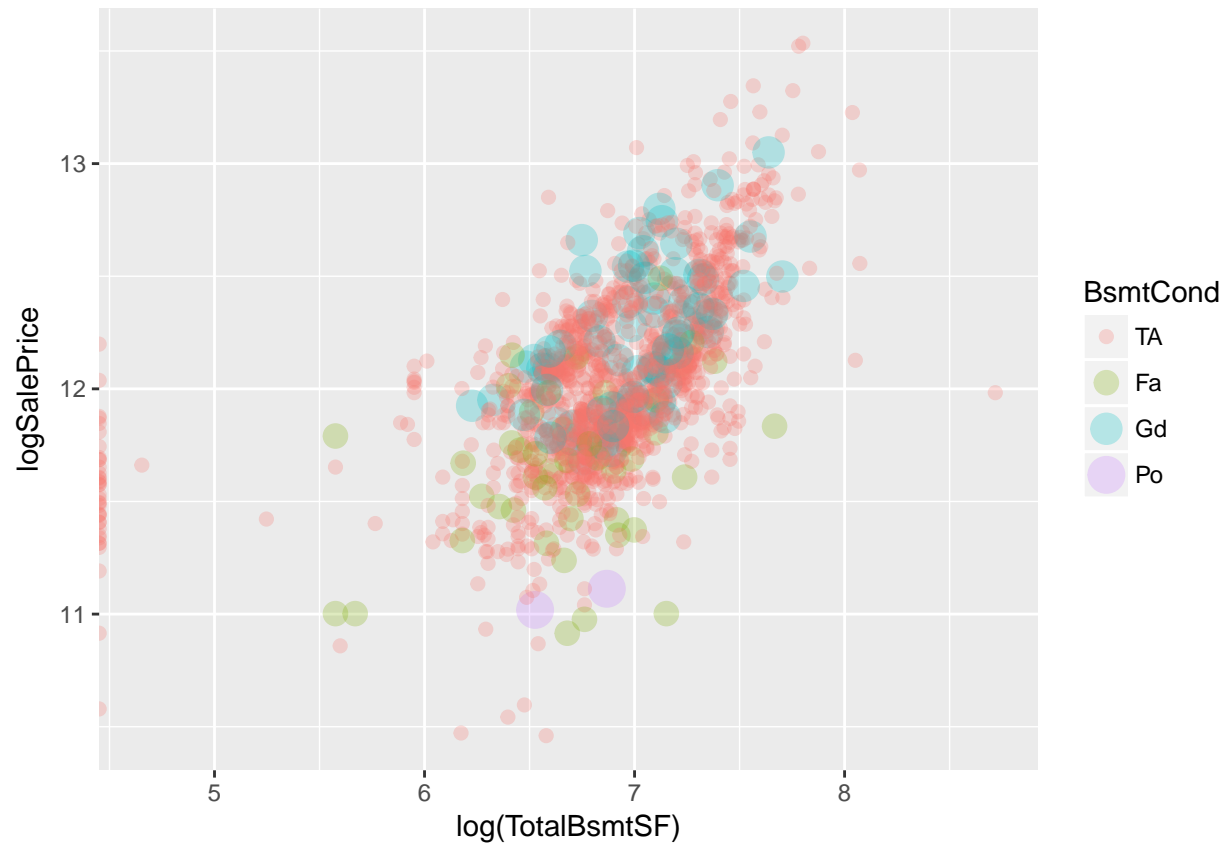
ggplot(train_na_filled) +
  geom_point(aes(x = log(LotArea), y = logSalePrice,
                 col = factor(OverallQual)),
             alpha = I(1/4))

```



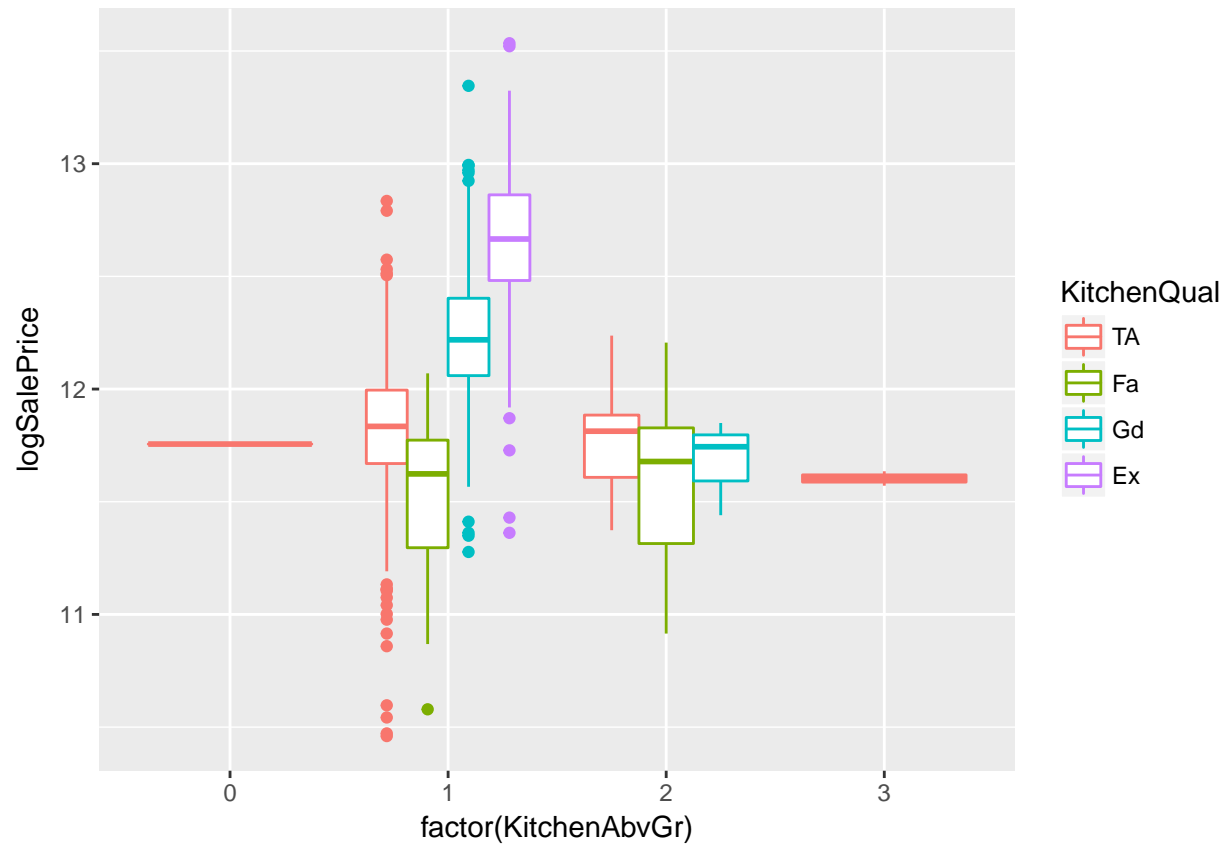
We can also get similar conclusion on the condition of basement, hence condition of basement is also a key factor count for housing price.

```
train_na_filled %>%
  mutate(BsmtCond = factor(BsmtCond, levels = c('TA', 'Fa', 'Gd', 'Po'))) %>%
  ggplot() +
  geom_point(aes(x = log(TotalBsmtSF), y = logSalePrice, col = BsmtCond,
                 size = BsmtCond), alpha = I(1/4))
```

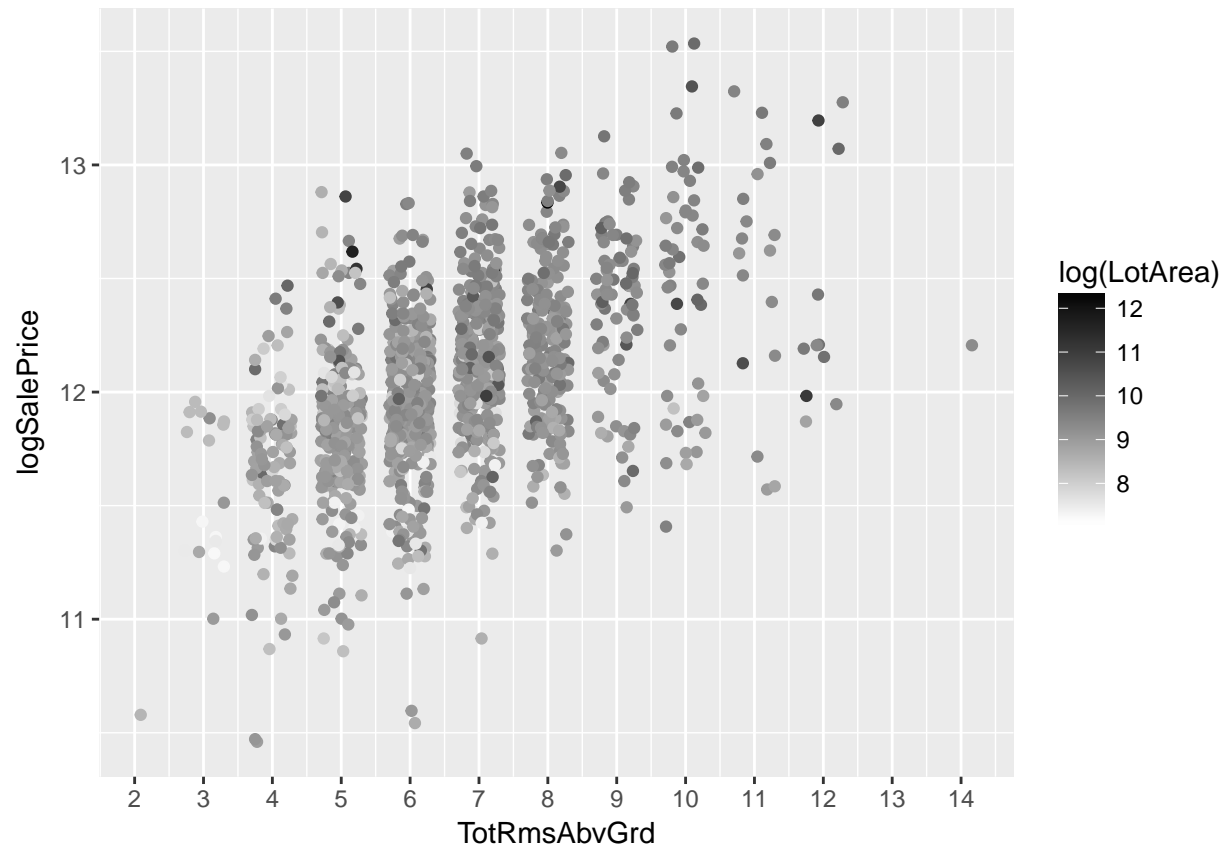


Then check the number and conditions of Bedrooms (above ground), Kitchen and overall rooms, we illustrate those variables along with *LotArea*.

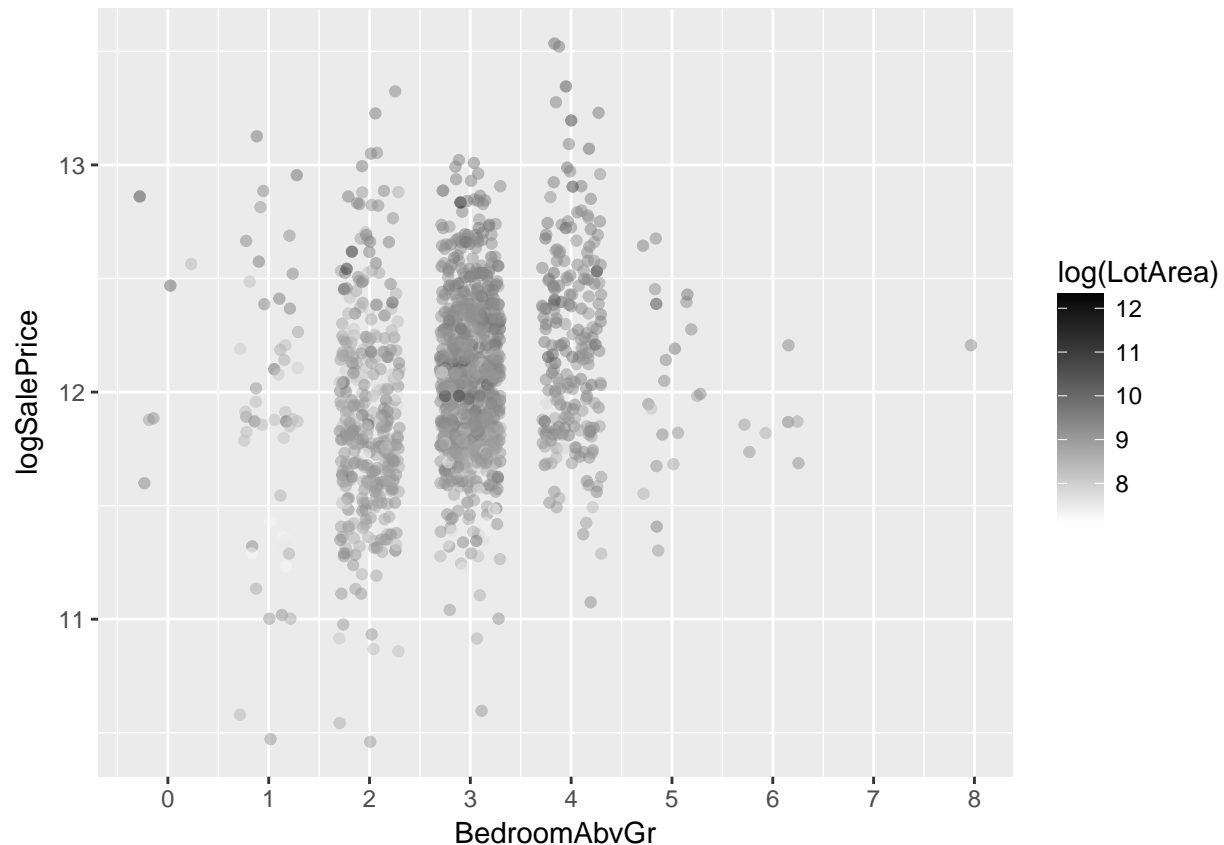
```
train_na_filled %>%
  mutate(KitchenQual = factor(KitchenQual, levels = c('TA', 'Fa', 'Gd', 'Ex'))) %>%
  ggplot(aes(x = factor(KitchenAbvGr), y = logSalePrice)) +
  geom_boxplot(aes(col = KitchenQual))
```



```
ggplot(train_na_filled, aes(x = TotRmsAbvGrd, y = logSalePrice)) +
  geom_jitter(aes(col = log(LotArea)), width = .3) +
  scale_x_continuous(breaks = 2:14) +
  scale_colour_gradient(low = "white", high = "black")
```



```
ggplot(train_na_filled) +
  geom_jitter(aes(x = BedroomAbvGr, y = logSalePrice, col = log(LotArea)),
    width = .3, alpha = I(1/2)) +
  scale_x_continuous(breaks = 0:8) +
  scale_colour_gradient(low = "white", high = "black")
```

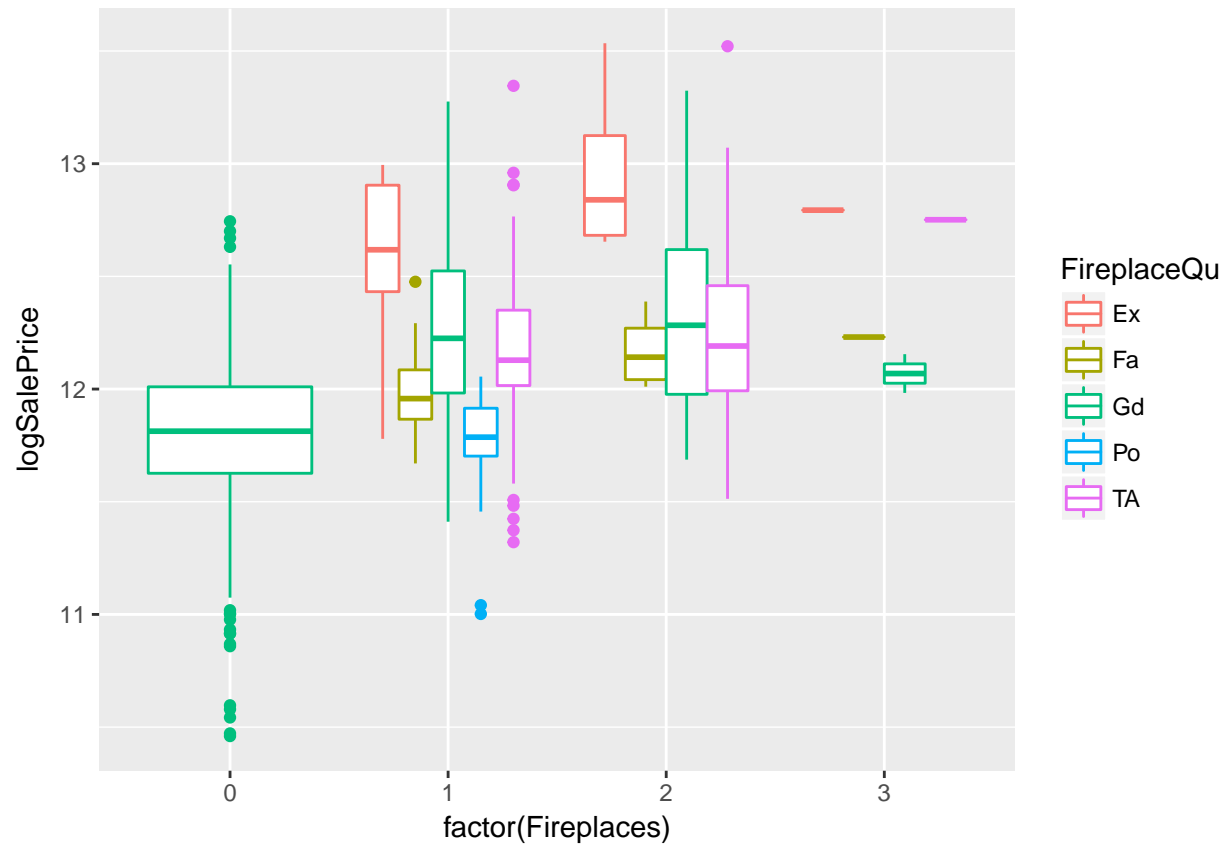


From the plots with *KitchenAbrGrd* and *KitchenQual* we find that many houses have just one kitchen, therefore the number of kitchen doesn't have too much contribution on *logSalePrice*. However *logSalePrice* shows significant differences on different levels of *KitchenQual*.

We see that *logSalePrice* increases as the total number of rooms increases in general, and similar pattern could be found on number of bedrooms but not very clear.

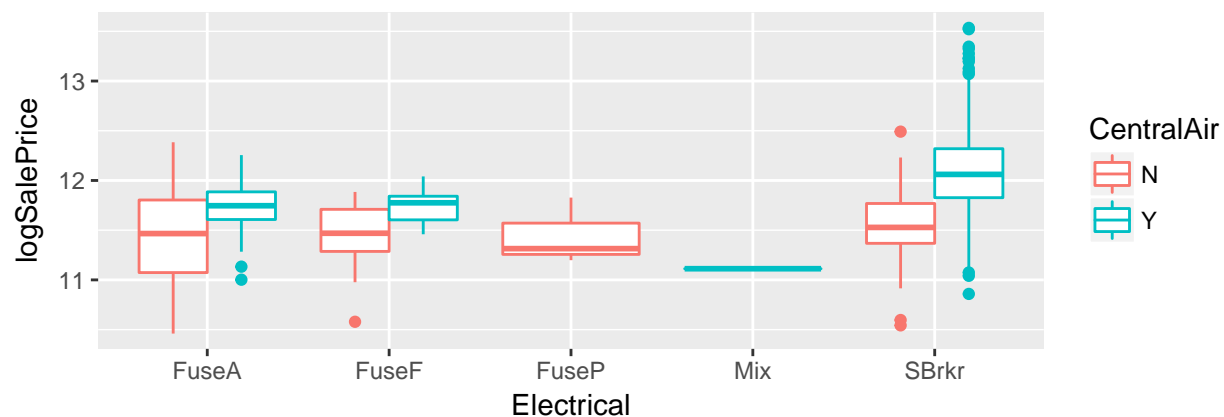
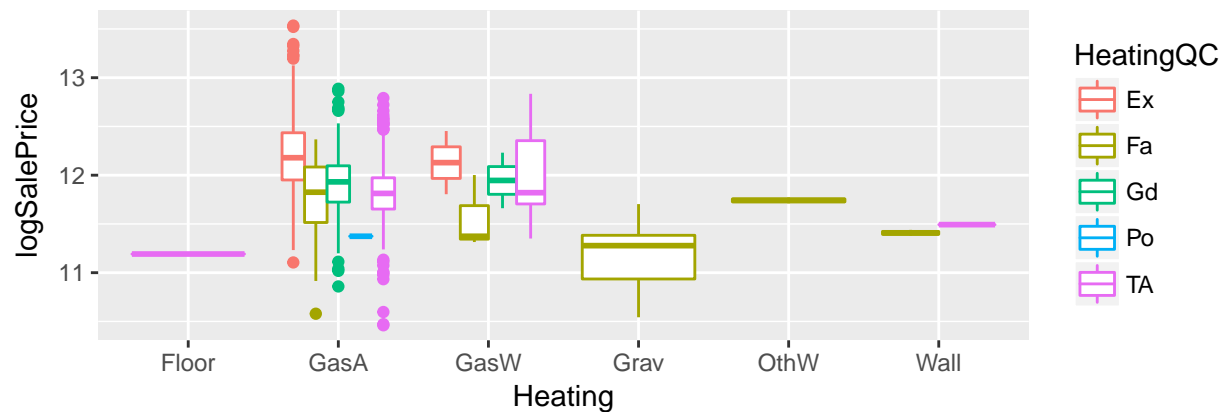
And what about *Fireplaces* and *FireplaceQu*?

```
ggplot(train_na_filled) +
  geom_boxplot(aes(x = factor(Fireplaces), y = logSalePrice,
    col = FireplaceQu))
```

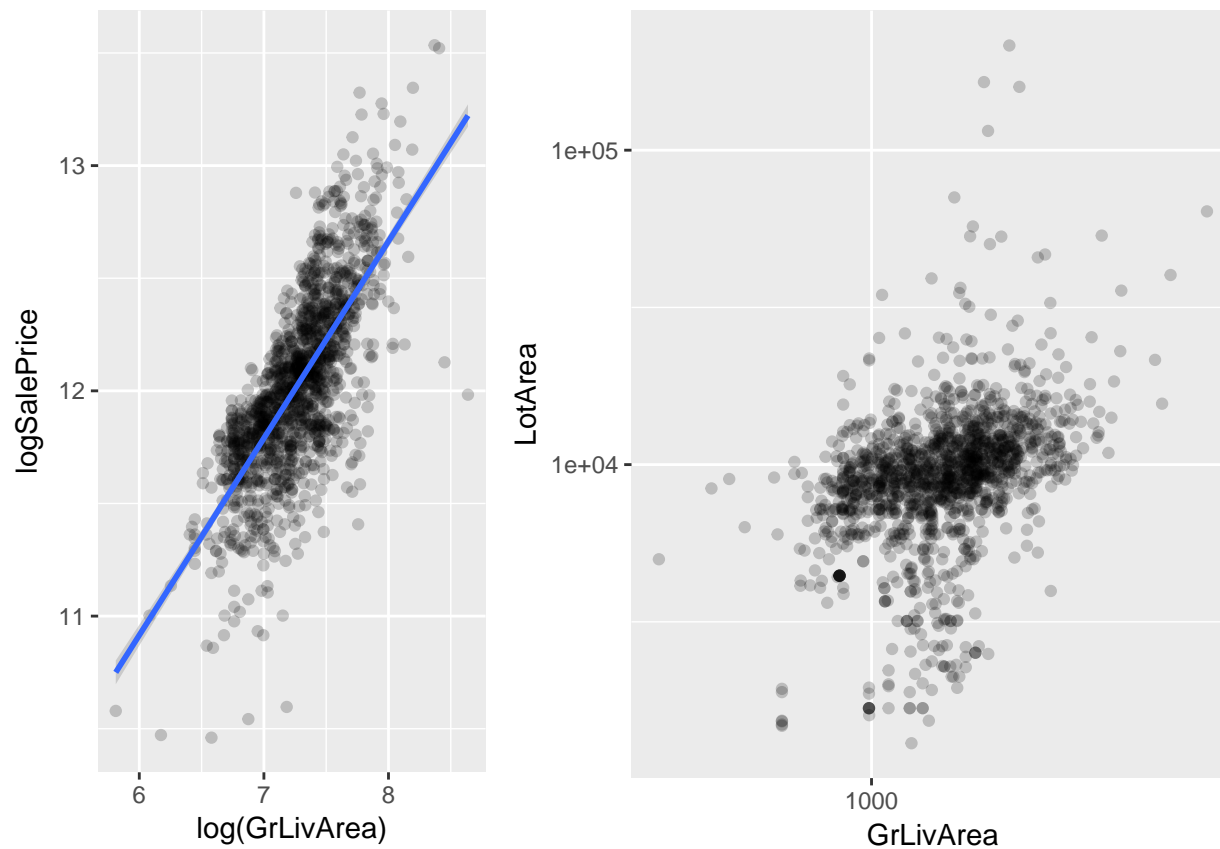


Electrical and *Heating* are also important factors in common sense, therefore we check them also as well as *GrLivArea*.

```
pHeat <- ggplot(train_na_filled) +
  geom_boxplot(aes(x = Heating, y = logSalePrice,
    col = HeatingQC))
pElecAir <- ggplot(train_na_filled) +
  geom_boxplot(aes(x = Electrical, y = logSalePrice,
    col = CentralAir))
grid.arrange(pHeat, pElecAir, nrow = 2)
```

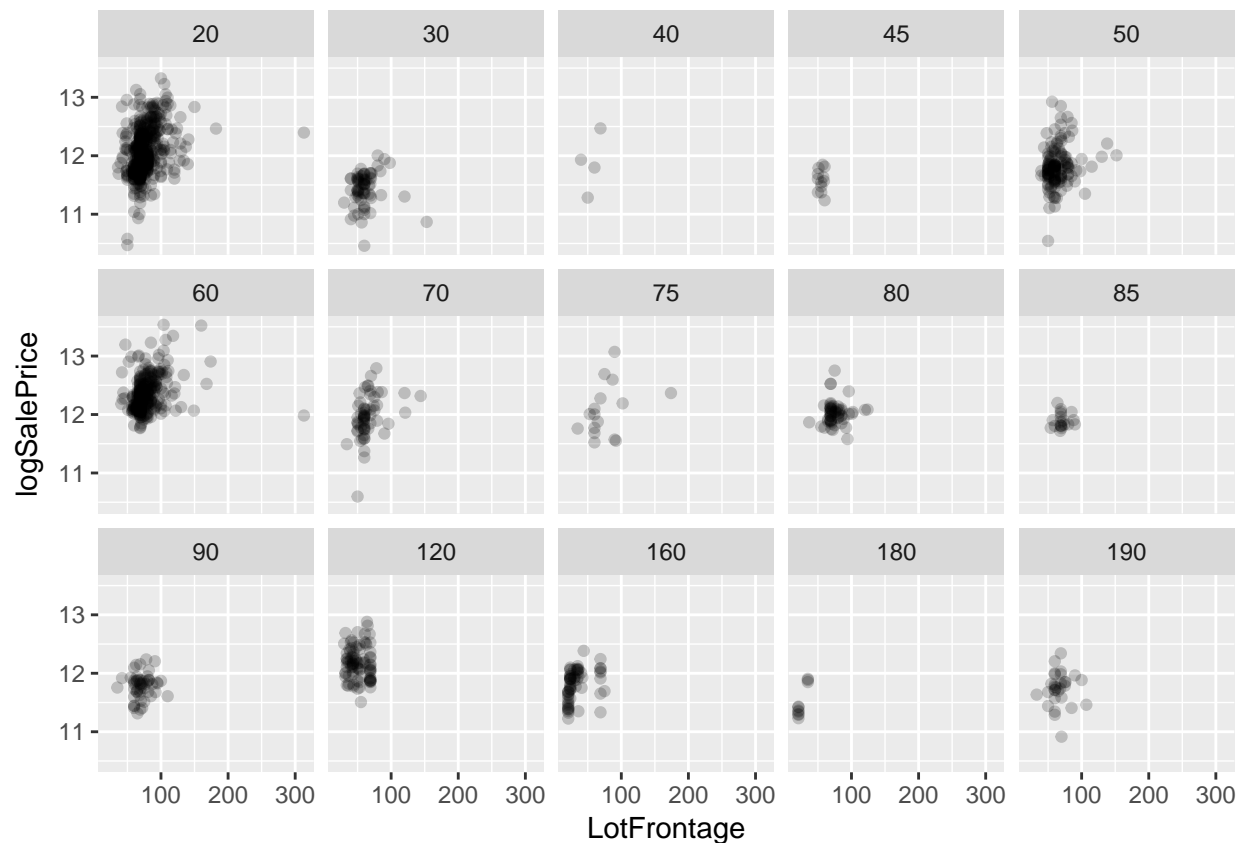
```
pGrLiv <- ggplot(train_na_filled, aes(x = log(GrLivArea), y = logSalePrice)) +
  geom_point(alpha = I(1/5)) +
  geom_smooth(method = 'lm', level = .9)
pGrLivLot <- ggplot(train_na_filled) +
  geom_point(aes(x = GrLivArea, y = LotArea), alpha = I(1/5)) +
  scale_x_log10() +
  scale_y_log10()
grid.arrange(pGrLiv, pGrLivLot, ncol = 2, widths = c(2, 3))
```



(4) Exterior

First we illustrate the relationships between various “Exterior” variables with *logSalePrice*,

```
ggplot(train_na_filled) +
  geom_point(aes(x = LotFrontage, y = logSalePrice), alpha = I(1/5)) +
  facet_wrap(~ MSSubClass, ncol = 5)
```



```

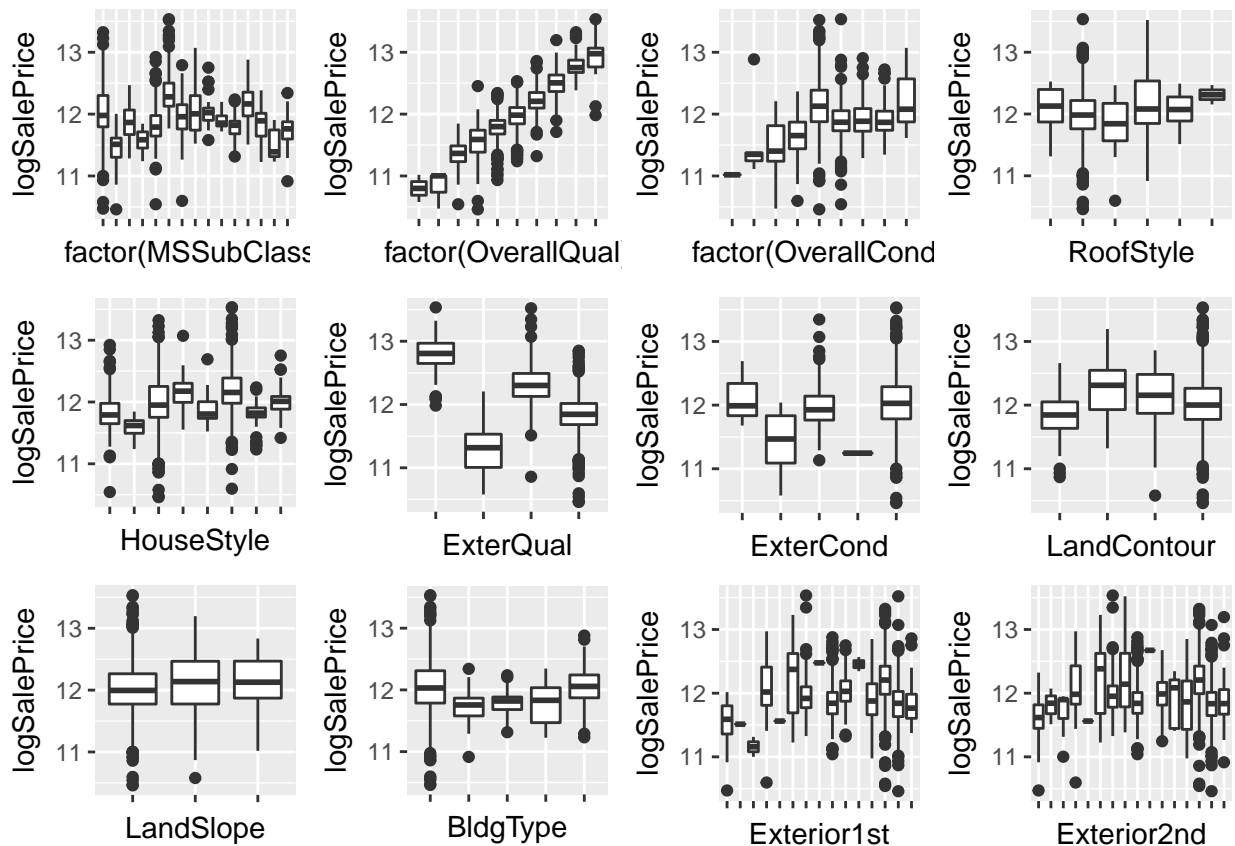
pMSSubClass <- ggplot(train_na_filled) +
  geom_boxplot(aes(x = factor(MSSubClass), y = logSalePrice)) +
  theme(axis.text.x = element_blank())
pOverallQu <- ggplot(train_na_filled) +
  geom_boxplot(aes(x = factor(OverallQual), y = logSalePrice)) +
  theme(axis.text.x = element_blank())
pOverallConv <- ggplot(train_na_filled) +
  geom_boxplot(aes(x = factor(OverallCond), y = logSalePrice)) +
  theme(axis.text.x = element_blank())
pRfStl <- ggplot(train_na_filled) +
  geom_boxplot(aes(x = RoofStyle, y = logSalePrice)) +
  theme(axis.text.x = element_blank())
pHsStl <- ggplot(train_na_filled) +
  geom_boxplot(aes(x = HouseStyle, y = logSalePrice)) +
  theme(axis.text.x = element_blank())
pExQu <- ggplot(train_na_filled) +
  geom_boxplot(aes(x = ExterQual, y = logSalePrice)) +
  theme(axis.text.x = element_blank())
pExCond <- ggplot(train_na_filled) +
  geom_boxplot(aes(x = ExterCond, y = logSalePrice)) +
  theme(axis.text.x = element_blank())
pEx1 <- ggplot(train_na_filled) +
  geom_boxplot(aes(x = Exterior1st, y = logSalePrice)) +
  theme(axis.text.x = element_blank())
pEx2 <- ggplot(train_na_filled) +
  geom_boxplot(aes(x = Exterior2nd, y = logSalePrice)) +

```

```

theme(axis.text.x = element_blank())
pLdConto <- ggplot(train_na_filled) +
  geom_boxplot(aes(x = LandContour, y = logSalePrice)) +
  theme(axis.text.x = element_blank())
pLdSlp <- ggplot(train_na_filled) +
  geom_boxplot(aes(x = LandSlope, y = logSalePrice)) +
  theme(axis.text.x = element_blank())
pBldgTp <- ggplot(train_na_filled) +
  geom_boxplot(aes(x = BldgType, y = logSalePrice)) +
  theme(axis.text.x = element_blank())
grid.arrange(pMSSubClass, pOverallQu, pOverallConv, pRfStl,
  pHsStl, pExQu, pExCond, pLdConto, pLdSlp, pBldgTp, pEx1,
  pEx2, ncol = 4)

```



Almost all boxplots show the differences of distribution of *logSalePrice* on different levels of “Exterior” variables. We can include all of them into models, but this might lead to collinearity.

(5) Location and transportation

Next we consider the location and transportation condition of house. Again, we use boxplots to display the data.

```

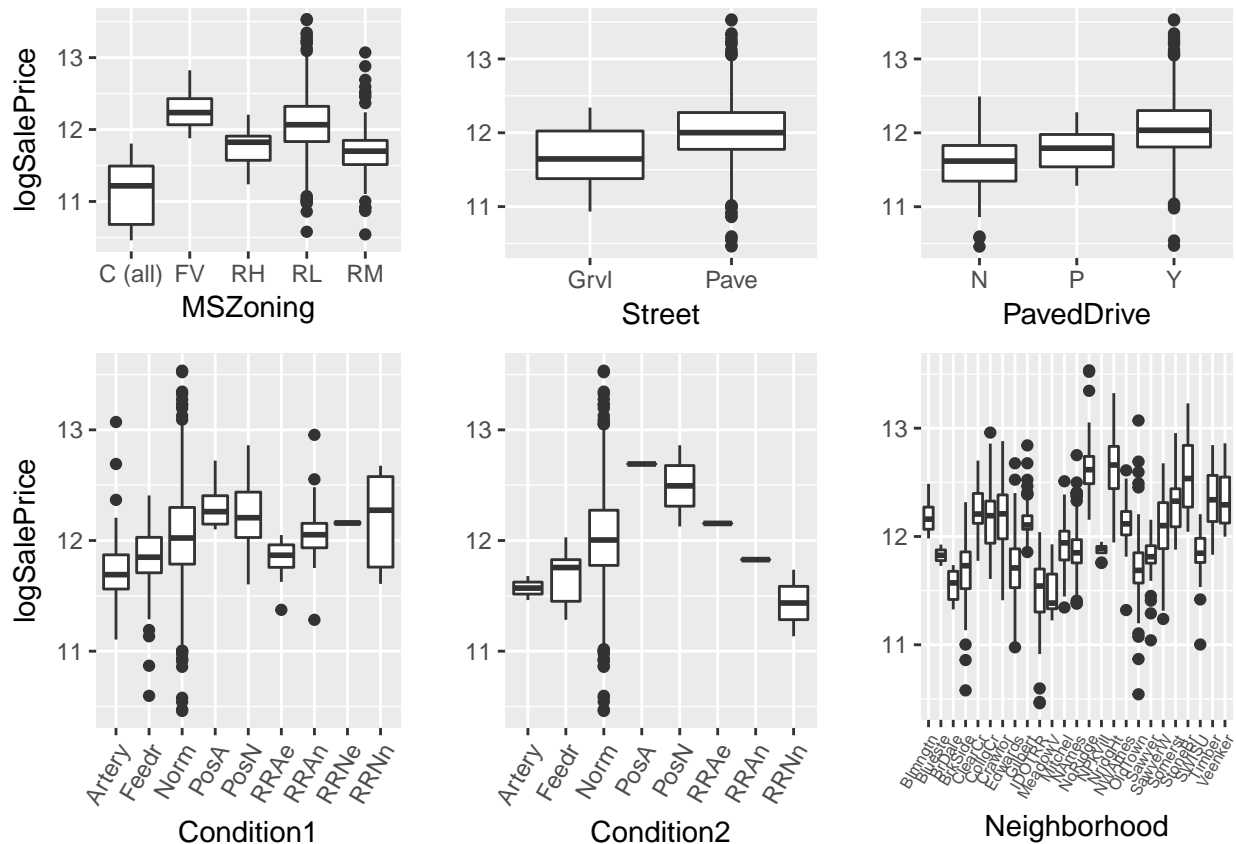
pMSzoning <- ggplot(train_na_filled) +
  geom_boxplot(aes(x = MSZoning, y = logSalePrice))
pNeighbor <- ggplot(train_na_filled) +
  geom_boxplot(aes(x = Neighborhood, y = logSalePrice)) +

```

```

theme(axis.text.x = element_text(angle = 60, hjust = 1, size = 5.5)) +
  ylab('')
pStreet <- ggplot(train_na_filled) +
  geom_boxplot(aes(x = Street, y = logSalePrice)) +
  ylab('')
pPDrive <- ggplot(train_na_filled) +
  geom_boxplot(aes(x = PavedDrive, y = logSalePrice)) +
  ylab('')
pCond1 <- ggplot(train_na_filled) +
  geom_boxplot(aes(x = Condition1, y = logSalePrice)) +
  theme(axis.text.x = element_text(angle = 60, hjust = 1))
pCond2 <- ggplot(train_na_filled) +
  geom_boxplot(aes(x = Condition2, y = logSalePrice)) +
  theme(axis.text.x = element_text(angle = 60, hjust = 1)) +
  ylab('')
grid.arrange(pMSzoning, pStreet, pPDrive,
              pCond1, pNeighbor, nrow = 2,
              heights = c(2, 3))

```



As a result we can claim that the location and transportation are highly correlate with the price of hoses.

Then check *Garage_*

```

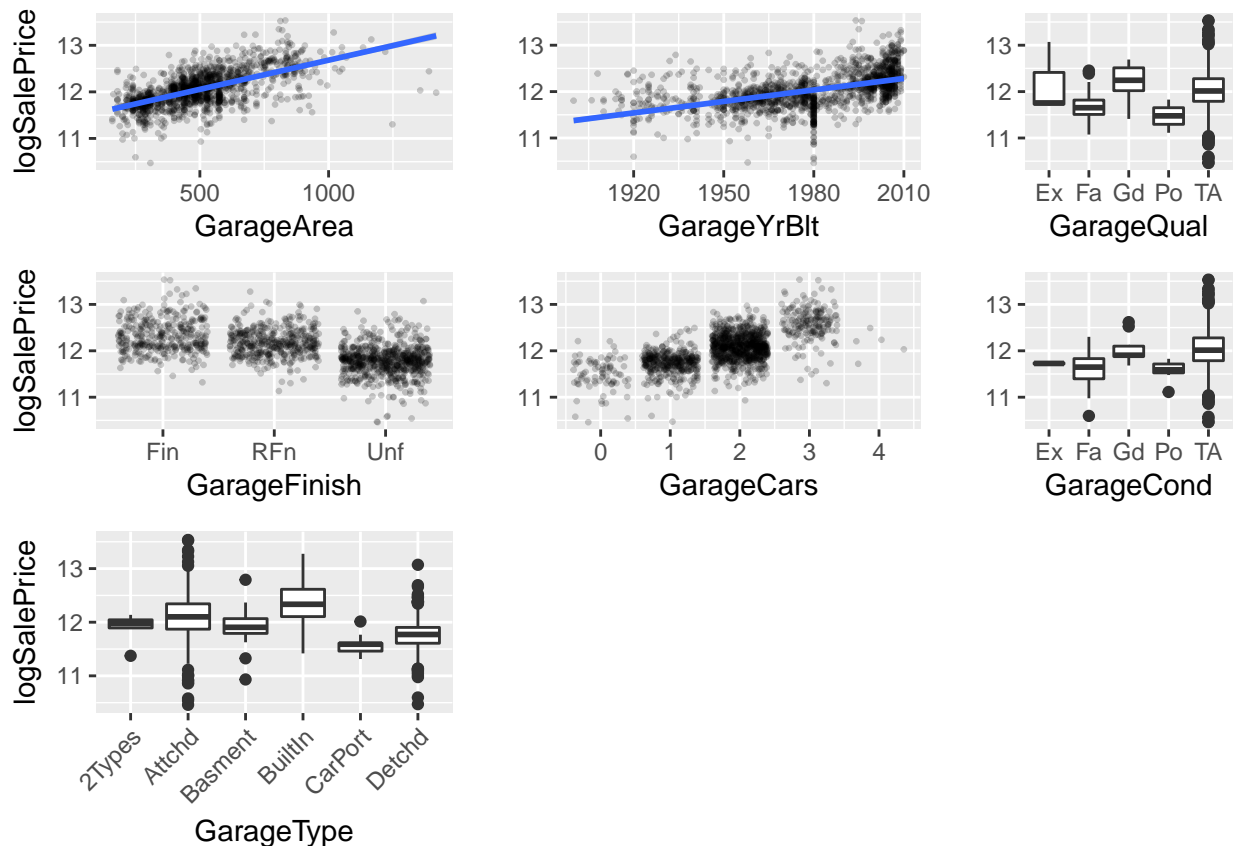
pGargTp <- ggplot(train_na_filled) +
  geom_boxplot(aes(x = GarageType, y = logSalePrice)) +
  theme(axis.text.x = element_text(angle = 45, hjust = 1))
pGargYBt <- ggplot(train_na_filled, aes(x = GarageYrBlt, y = logSalePrice)) +

```

```

  geom_point(alpha = I(1/5), size = .5) +
  geom_smooth(method = 'lm') +
  ylab('')
pGargFns <- ggplot(train_na_filled, aes(x = GarageFinish, y = logSalePrice)) +
  geom_jitter(alpha = I(1/5), size = .5)
pGargCars <- ggplot(train_na_filled, aes(x = GarageCars, y = logSalePrice)) +
  geom_jitter(alpha = I(1/5), size = .5) +
  ylab('')
pGargAr <- ggplot(train_na_filled %>% filter(GarageArea > 0),
  aes(x = GarageArea, y = logSalePrice)) +
  geom_point(alpha = I(1/5), size = .5) +
  geom_smooth(method = 'lm')
pGargQu <- ggplot(train_na_filled, aes(x = GarageQual, y = logSalePrice)) +
  geom_boxplot() +
  ylab('')
pGargCond <- ggplot(train_na_filled, aes(x = GarageCond, y = logSalePrice)) +
  geom_boxplot() +
  ylab('')
grid.arrange(pGargAr, pGargYBt, pGargQu,
  pGargFns, pGargCars,
  pGargCond, pGargTp,
  nrow = 3, heights = c(3, 3, 4),
  widths = c(3, 3, 2))

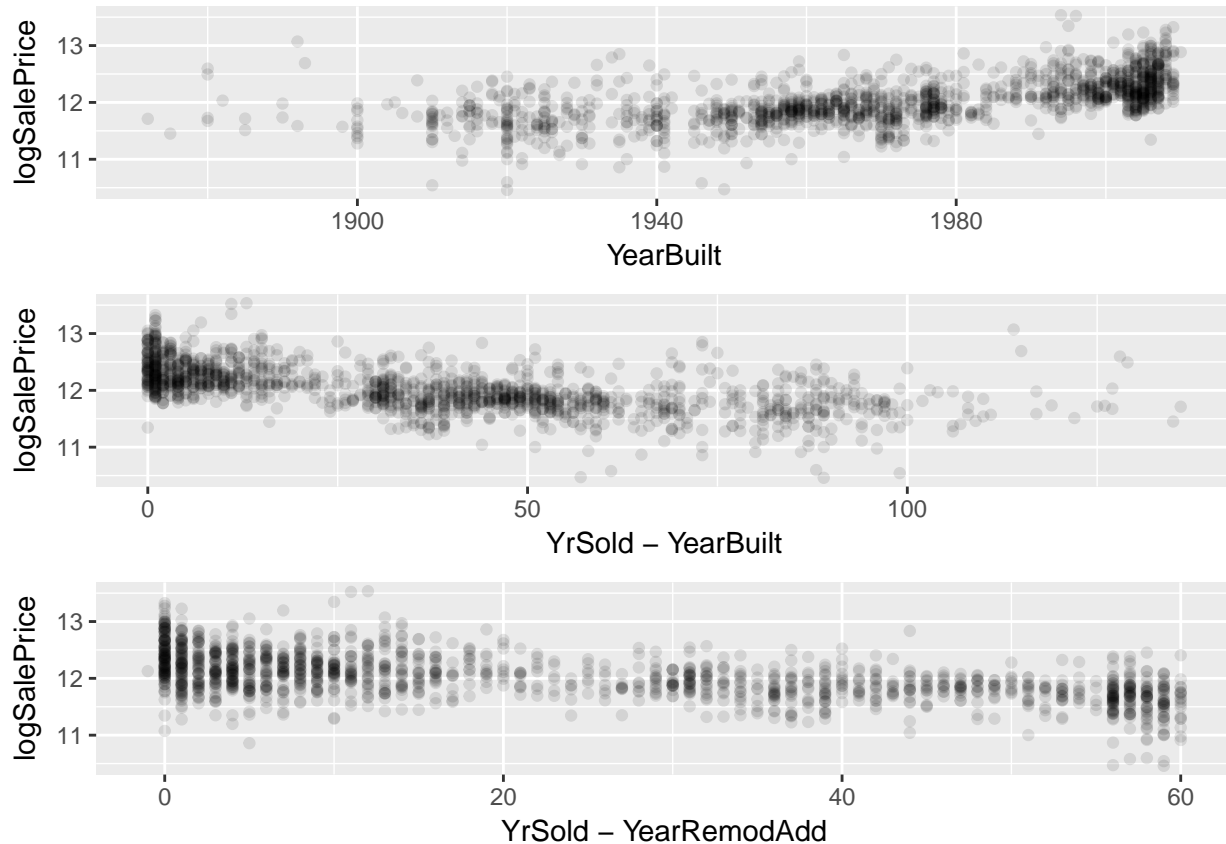
```



We can see that the distribution of *logSalePrice* is similar on *GarageQual* and *GarageCond*. And generally speaking, *GarageArea* should have a high positive correlation with *GarageCars*.

(6) Other attributes

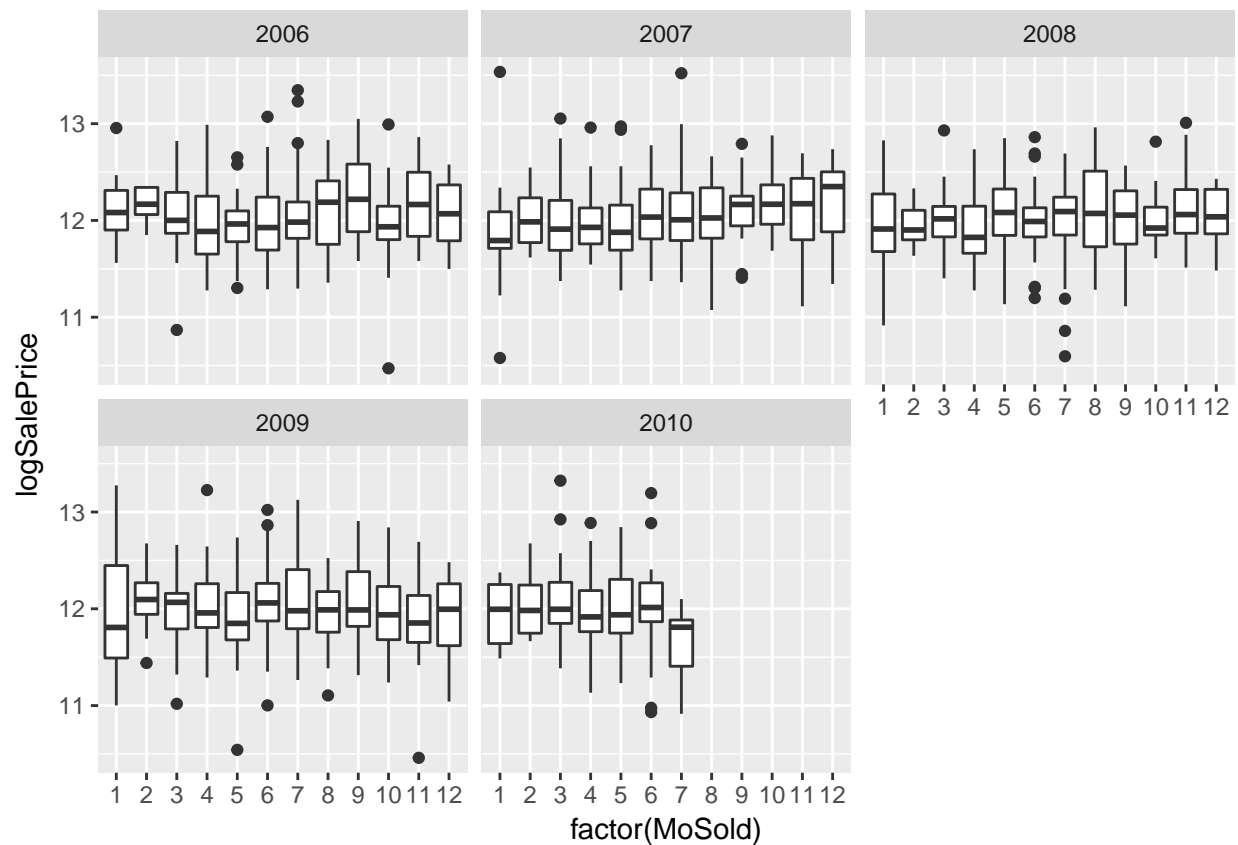
```
pYrBlt <- ggplot(train_na_filled) +  
  geom_point(aes(x = YearBuilt, y = logSalePrice), alpha = I(1/10))  
pHsAge <- ggplot(train_na_filled) +  
  geom_point(aes(x = YrSold - YearBuilt, y = logSalePrice), alpha = I(1/10))  
pRemd <- ggplot(train_na_filled, aes(x = YrSold - YearRemodAdd, y = logSalePrice)) +  
  geom_point(alpha = I(1/10))  
grid.arrange(pYrBlt, pHsAge, pRemd, nrow = 3, heights = c(2, 2, 2))
```



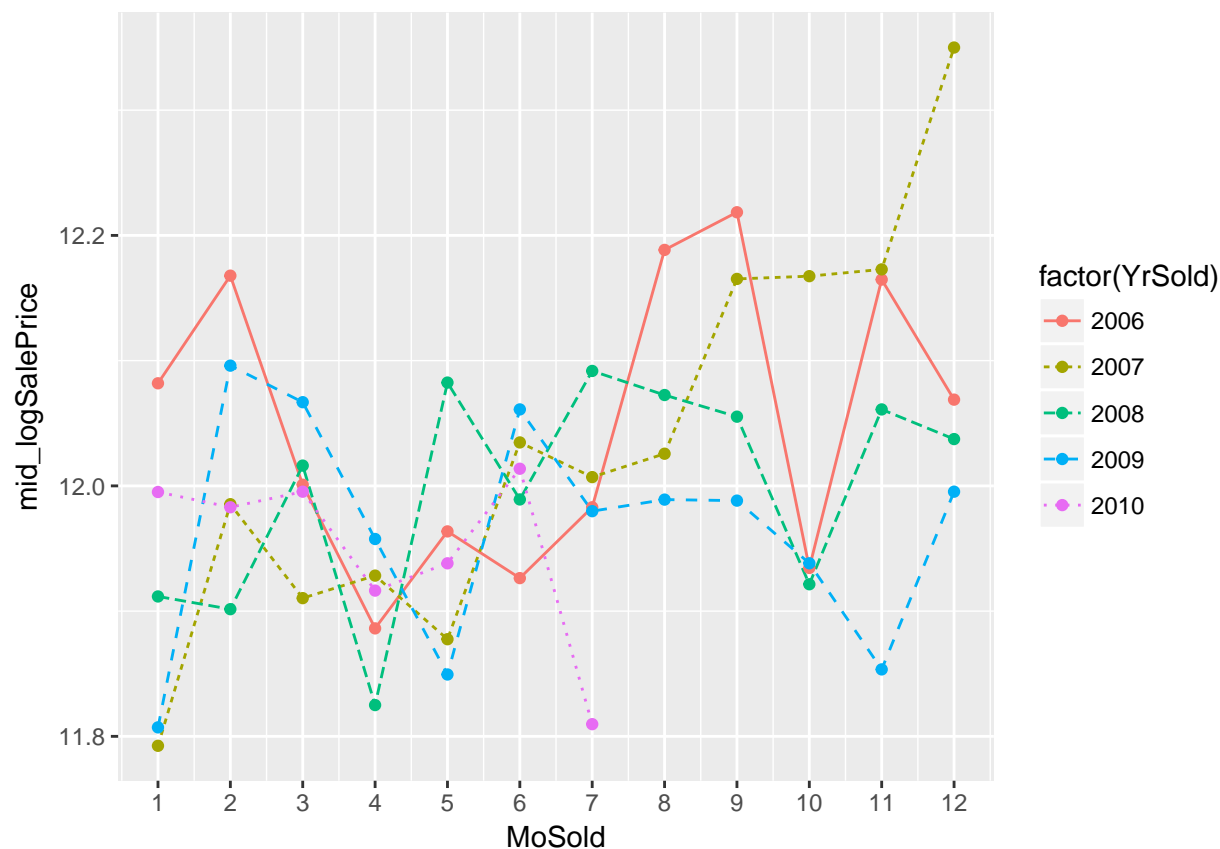
The scatter plots between *Year* attributes and *logSalePrice* express that the “Age” of house has linear correlation with its value in general, the variation of price of houses built prior to 1900 is higher.

Then we will check *MoSold* and *YrSold* to see if they related with the price of houses. We bet that it might exists seasonal trend

```
ggplot(train_na_filled) +  
  geom_boxplot(aes(x = factor(MoSold), y = logSalePrice)) +  
  facet_wrap(~ YrSold)
```



```
train_na_filled %>%
  group_by(MoSold, YrSold) %>%
  summarise(mid_logSalePrice = median(logSalePrice)) %>%
  ggplot(aes(x = MoSold, y = mid_logSalePrice)) +
  geom_line(aes(group = YrSold,
                col = factor(YrSold),
                linetype = factor(YrSold))) +
  geom_point(aes(group = YrSold,
                 col = factor(YrSold),
                 linetype = factor(YrSold))) +
  scale_x_continuous(breaks = 1:12)
```

Also we check the relationship between *SaleType*, *SaleCondisiton* and *MoSold*.

```
options(digits = 2)
# prop of different SaleType above SaleCondition
table(train_na_filled$SaleType, train_na_filled$SaleCondition) /
  rowSums(table(train_na_filled$SaleType, train_na_filled$SaleCondition))
```

```
##
##      Abnorml AdjLand Alloca Family Normal Partial
## COD      0.5581  0.0000  0.0000  0.0000  0.4419  0.0000
## Con       0.0000  0.0000  0.0000  0.0000  1.0000  0.0000
## ConLD     0.2222  0.0000  0.0000  0.0000  0.6667  0.1111
## ConLI     0.2000  0.0000  0.0000  0.0000  0.8000  0.0000
## ConLw     0.0000  0.0000  0.0000  0.0000  1.0000  0.0000
## CWD       0.2500  0.0000  0.0000  0.2500  0.5000  0.0000
## New       0.0000  0.0000  0.0000  0.0000  0.0000  1.0000
## Oth       1.0000  0.0000  0.0000  0.0000  0.0000  0.0000
## WD        0.0552  0.0032  0.0095  0.0150  0.9155  0.0016
```

```
# prop of different MoSold above SaleCondition
table(train_na_filled$MoSold, train_na_filled$SaleCondition) /
  rowSums(table(train_na_filled$MoSold, train_na_filled$SaleCondition))
```

```
##
##      Abnorml AdjLand Alloca Family Normal Partial
## 1      0.1207  0.0000  0.0172  0.0345  0.7414  0.0862
## 2      0.0769  0.0000  0.0000  0.0000  0.8846  0.0385
## 3      0.0755  0.0000  0.0283  0.0189  0.8113  0.0660
```

```
## 4 0.0780 0.0071 0.0000 0.0142 0.8369 0.0638
## 5 0.0588 0.0049 0.0000 0.0049 0.8824 0.0490
## 6 0.0237 0.0040 0.0198 0.0119 0.8814 0.0593
## 7 0.0684 0.0000 0.0000 0.0085 0.8291 0.0940
## 8 0.0984 0.0000 0.0082 0.0246 0.7377 0.1311
## 9 0.0476 0.0000 0.0159 0.0317 0.7302 0.1746
## 10 0.1011 0.0112 0.0000 0.0112 0.7528 0.1236
## 11 0.0886 0.0000 0.0127 0.0127 0.7595 0.1266
## 12 0.1017 0.0000 0.0000 0.0169 0.7627 0.1186
```

Saving data.

```
save.image("C:/Users/Bangda/Desktop/kaggle/housing-price/eda.RData")
```